

ETP story-formalization benchmark

Stories generated from Equational Theories Project implications (`storyform.py`), formalized back by each model over OpenRouter, and graded syntactically (`checkform.py`). Correct = exact, or an accepted symmetry (sides swapped, or both equations uniformly dualized).


results/run-s0-n30-think-off · results/run-s0-n30-think-on · results/run-strat5-s0-think-off · results/run-strat5-s0-think-on

Runs	Models	Stories	Graded calls	Overall correct
4	8	70	1,120	67%

30 uniformly sampled pairs · seed 0 · no-think vs all-think

FIGURE 1 Correct rate by model: no-think vs all-think

Same 30 stories under both regimes; each row connects a model's cold-pass rate to its deliberate rate.




► Data: correct rate per model and regime

FIGURE 2 **Verdict composition — no-think regime**

How each model's answers grade out; the right-hand figure is the correct rate (exact + swapped + dualized).

■ exact ■ correct-swapped ■ correct-dualized ■ wrong ■ unparseable




► Data: verdict counts per model

FIGURE 3 **Verdict composition — all-think regime**

How each model's answers grade out; the right-hand figure is the correct rate (exact + swapped + dualized).

■ exact ■ correct-swapped ■ correct-dualized ■ wrong ■ unparseable




► Data: verdict counts per model

40 pairs, 5 per operation-count bin 1–8 · seed 0 · no-think vs all-think

FIGURE 4 **Correct rate by model: no-think vs all-think**


Same 40 stories under both regimes; each row connects a model's cold-pass rate to its deliberate rate.



► Data: correct rate per model and regime

FIGURE 5 **Accuracy vs implication complexity**

Correct rate pooled over all models, by the implication's total operation count (both equations combined).




► Data: pooled correct rate per operation-count bin

FIGURE 6 **Verdict composition — no-think regime**

How each model's answers grade out; the right-hand figure is the correct rate (exact + swapped + dualized).

■ exact ■ correct-swapped ■ correct-dualized ■ wrong ■ unparseable




► Data: verdict counts per model

FIGURE 7 **Verdict composition — all-think regime**

How each model's answers grade out; the right-hand figure is the correct rate (exact + swapped + dualized).

■ exact ■ correct-swapped ■ correct-dualized ■ wrong ■ unparseable



► Data: verdict counts per model